I started Lab 6 by comparing tokenizers on six inputs, and I saw that the regex splits can't into can and t and breaks emails and version numbers, while whitespace keeps them whole, so I kept the original text for reporting. Then I wrote six documents of 150 to 250 words each about late submission, passwords, assignments, remote work and printers, and I chunked them with 40-word windows with overlap 0 and 10 and with sentence-aware packing, with source offsets saved. The hardest part was the model download, and my sentence splitter cut wrongly on abbreviations like Dr., so I checked the chunks by hand. For retrieval I ran all three chunkings with TF-IDF and with real MiniLM embeddings on eight answerable and two unanswerable queries. Overlap 10 with embeddings was best: Hit@3 1.000 and MRR@3 1.000 on development queries and 1.000 and 0.750 on the test set, while TF-IDF fell to 0.250 because the test queries were paraphrases. I also learned that a high score does not mean a right answer, because the flight-prices query still got a TF-IDF score of 0.345 even though no document could answer it.